Day 1 - Recap

→ Generative AI → 'FMs | LLMs'
→ Amazon Bedrock → "250 + FMs"
APIs → InvokeModel
→ Converse

→ FM / LLM → 'Inference' Parameters / Configurations

0-1 / 0-2
- Temperature ✓ More focused / More creative -
- Top 'P' (0-1) 0.9 (90%) 0.95 (95%)
- Top K ('Claude' family) [250]
- max-tokens (Output Response) - [2K]

'Inference Profile'

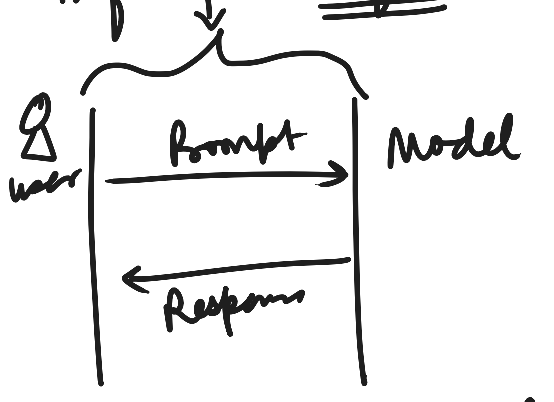

When ?
1. Global application
2. HA
3. Prod system

Model ID : us. or global. amazon . Nova - lite - v.1 ←

Why ?
1. Cost allocation Tags
2. usage tracking for app
3. Seperate env based
   - prod-ai-profile
   - dev-ai-profile
   - fin-genai-profile.

Day 2 - Agenda

* 1. RAG → Amazon 'Bedrock' KB
* 2. Opensource Frameworks → Langchain
3. Evaluate Gen AI Application Component
   → 'RAGAS' → Lab
   * → Model Response evaluation → Demo.
4. Guardrails the GenAI App
   ↳ Amazon Bedrock Guardrails. "Lab"
5. Agents, Tools, (ReACT)
   Amazon Bedrock Agent
6. Agentic AI, Agentcore → Concept

Lab - Capstone Project

RAG ?

Retrieval ⇒ user Input → search an External Info / Knowledge.
Augmentation ⇒ Add that retrieved info to the user Input
Generation ⇒ LLM uses the Augmented prompt to produce a response

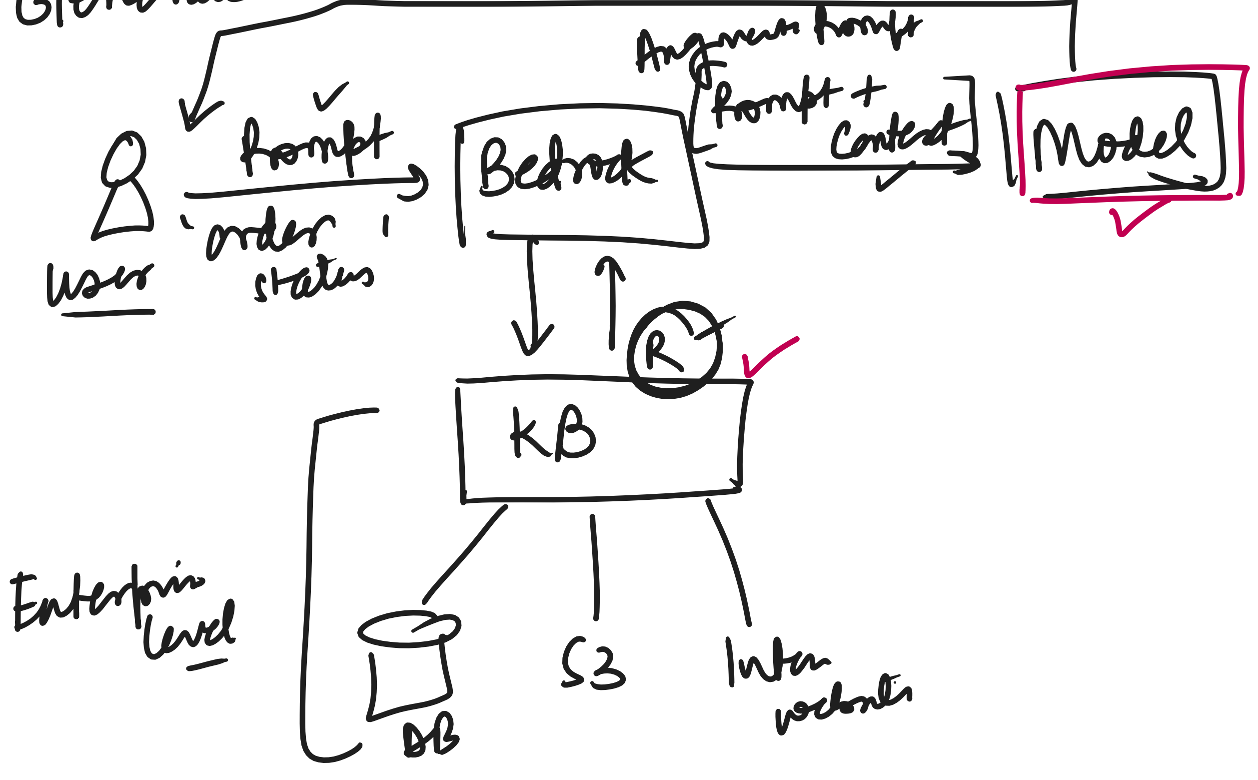

Bedrock APIs	Open Source
More control	More Abstraction
AWS SDK updates	Community update
Efficient API	Faster App Composition
Good for simple	Good for chains, memory, RAG, Agents

"Responsible AI" Dimensions

1. Controllability
2. Privacy & Security
3. Safety
4. Fairness
5. Veracity / Robustness
6. Explainability
7. Transparency
8. Governance.